



Choosing the right model

Selecting the optimal Claude model for your application involves balancing three key considerations: capabilities, speed, and cost. This guide helps you make an informed decision based on your specific requirements.

[Copy page](#)

Establish key criteria

When choosing a Claude model, we recommend first evaluating these factors:

- **Capabilities:** What specific features or capabilities will you need the model to have in order to meet your needs?
- **Speed:** How quickly does the model need to respond in your application?
- **Cost:** What's your budget for both development and production usage?

Knowing these answers in advance will make narrowing down and deciding which model to use much easier.

Choose the best model to start with

There are two general approaches you can use to start testing which Claude model best works for your needs.

Option 1: Start with a fast, cost-effective model

For many applications, starting with a faster, more cost-effective model like Claude Haiku 4.5 can be the optimal approach:

1. Begin implementation with Claude Haiku 4.5
2. Test your use case thoroughly
3. Evaluate if performance meets your requirements

[Ask Docs](#)

This approach allows for quick iteration, lower development costs, and is often sufficient for many common applications. This approach is best for:

- Initial prototyping and development
- Applications with tight latency requirements
- Cost-sensitive implementations
- High-volume, straightforward tasks

Option 2: Start with the most capable model

For complex tasks where intelligence and advanced capabilities are paramount, you may want to start with the most capable model and then consider optimizing to more efficient models down the line:

1. Implement with Claude Sonnet 4.5
2. Optimize your prompts for these models
3. Evaluate if performance meets your requirements
4. Consider increasing efficiency by downgrading intelligence over time with greater workflow optimization

This approach is best for:

- Complex reasoning tasks
- Scientific or mathematical applications
- Tasks requiring nuanced understanding
- Applications where accuracy outweighs cost considerations
- Advanced coding

Model selection matrix

When you need...	We recommend starting with...	Example use cases
Best model for complex agents and coding, highest intelligence across most tasks, superior tool orchestration for long-running autonomous tasks	Claude Sonnet 4.5	Autonomous coding agents, cybersecurity automation, complex financial analysis, multi-hour research tasks, multi agent frameworks
Maximum intelligence with practical performance for complex specialized tasks	Claude Opus 4.5	Professional software engineering, advanced agents for office tasks, computer and

		browser use at scale, step-change vision applications
Exceptional intelligence and reasoning for specialized complex tasks	Claude Opus 4.1	Highly complex codebase refactoring, nuanced creative writing, specialized scientific analysis
Near-frontier performance with lightning-fast speed and extended thinking - our fastest and most intelligent Haiku model at the most economical price point	Claude Haiku 4.5	Real-time applications, high-volume intelligent processing, cost-sensitive deployments needing strong reasoning, sub-agent tasks

Decide whether to upgrade or change models

To determine if you need to upgrade or change models, you should:

1. Create benchmark tests specific to your use case - having a good evaluation set is the most important step in the process
2. Test with your actual prompts and data
3. Compare performance across models for:
 - Accuracy of responses
 - Response quality
 - Handling of edge cases
4. Weigh performance and cost tradeoffs

Next steps



Model comparison chart

See detailed specifications and pricing for the latest Claude models



What's new in Claude 4.5

Explore the latest improvements in Claude 4.5 models



 Claude Docs



Solutions	Learn	Help and security
AI agents	Blog	Availability
Code modernization	Catalog	Status
Coding	Courses	Support
Customer support	Use cases	Discord
Education	Connectors	
Financial services	Customer stories	Terms and policies
Government	Engineering at Anthropic	Privacy policy
Life sciences	Events	Responsible disclosure policy
	Powered by Claude	Terms of service:
Partners	Service partners	Commercial
Amazon Bedrock	Startups program	Terms of service:
Google Cloud's Vertex AI		Consumer
	Company	Usage policy
	Anthropic	
	Careers	
	Economic Futures	
	Research	
	News	
	Responsible Scaling Policy	
	Security and compliance	
	Transparency	