

API Overview



 Copy page 

The Claude API is a RESTful API at <https://api.anthropic.com> that provides programmatic access to Claude models. The primary API is the Messages API (`POST /v1/messages`) for conversational interactions.

 **New to Claude?** Start with [Get started](#) for prerequisites and your first API call, or see [Working with Messages](#) for request/response patterns and examples.

Prerequisites

To use the Claude API, you'll need:

- An [Anthropic Console account](#)
- An [API key](#)

For step-by-step setup instructions, see [Get started](#).

Available APIs

The Claude API includes the following APIs:

General Availability:

- **[Messages API](#):** Send messages to Claude for conversational interactions (`POST /v1/messages`)
- **[Message Batches API](#):** Process large volumes of Messages requests asynchronously with 50% cost reduction (`POST /v1/messages/batches`)
- **[Token Counting API](#):** Count tokens in a message before sending to manage costs and rate limits (`POST /v1/messages/count_tokens`)
- **[Models API](#):** List available Claude models and their details (`GET /v1/models`)

Beta:

Ask Docs

- **Skills API:** Create and manage custom agent skills (`POST /v1/skills` , `GET /v1/skills`)

For the complete API reference with all endpoints, parameters, and response schemas, explore the API reference pages listed in the navigation. To access beta features, see [Beta headers](#).

Authentication

All requests to the Claude API must include these headers:

Header	Value	Required
x-api-key	Your API key from Console	Yes
anthropic-version	API version (e.g., <code>2023-06-01</code>)	Yes
content-type	<code>application/json</code>	Yes

If you are using the [Client SDKs](#), the SDK will send these headers automatically. For API versioning details, see [API versions](#).

Getting API Keys

The API is made available via the web [Console](#). You can use the [Workbench](#) to try out the API in the browser and then generate API keys in [Account Settings](#). Use [workspaces](#) to segment your API keys and [control spend](#) by use case.

Client SDKs

Anthropic provides official SDKs that simplify API integration by handling authentication, request formatting, error handling, and more.

Benefits:

- Automatic header management (x-api-key, anthropic-version, content-type)
- Type-safe request and response handling
- Built-in retry logic and error handling
- Streaming support
- Request timeouts and connection management

Example (Python):

```
client = Anthropic() # Reads ANTHROPIC_API_KEY from environment
message = client.messages.create(
    model="claude-sonnet-4-5",
    max_tokens=1024,
    messages=[{"role": "user", "content": "Hello, Claude"}]
)
```

For a list of client SDKs and their respective installation instructions, see [Client SDKs](#).

Claude API vs Third-Party Platforms

Claude is available through Anthropic's direct API and through partner platforms. Choose based on your infrastructure, compliance requirements, and pricing preferences.

Claude API

- **Direct access** to the latest models and features first
- **Anthropic billing and support**
- **Best for:** New integrations, full feature access, direct relationship with Anthropic

Third-Party Platform APIs

Access Claude through AWS, Google Cloud, or Microsoft Azure:

- **Integrated** with cloud provider billing and IAM
- **May have feature delays** or differences from the direct API
- **Best for:** Existing cloud commitments, specific compliance requirements, consolidated cloud billing

Platform	Provider	Documentation
Amazon Bedrock	AWS	Claude on Amazon Bedrock
Vertex AI	Google Cloud	Claude on Vertex AI
Azure AI	Microsoft Azure	Claude on Azure AI

Request Size Limits

The API has different maximum request sizes depending on the endpoint:

Endpoint	Maximum Size
Standard endpoints (Messages, Token Counting)	32 MB
Batch API	256 MB
Files API	500 MB

If you exceed these limits, you'll receive a 413 `request_too_large` error.

Response Headers

The Claude API includes the following headers in every response:

- `request-id` : A globally unique identifier for the request
- `anthropic-organization-id` : The organization ID associated with the API key used in the request

Rate Limits and Availability

Rate Limits

The API enforces rate limits and spend limits to prevent misuse and manage capacity. Limits are organized into usage tiers that increase automatically as you use the API. Each tier has:

- Spend limits:** Maximum monthly cost for API usage
- Rate limits:** Maximum number of requests per minute (RPM) and tokens per minute (TPM)

You can view your organization's current limits in the [Console](#). For higher limits or Priority Tier (enhanced service levels with committed spend), contact sales through the Console.

For detailed information about limits, tiers, and the token bucket algorithm used for rate limiting, see [Rate limits](#).

Availability

The Claude API is available in [many countries and regions](#) worldwide. Check the supported regions page to confirm availability in your location.

Here's a minimal request using the Messages API:

```
curl https://api.anthropic.com/v1/messages \
--header "x-api-key: $ANTHROPIC_API_KEY" \
--header "anthropic-version: 2023-06-01" \
--header "content-type: application/json" \
--data '{
  "model": "claude-sonnet-4-5",
  "max_tokens": 1024,
  "messages": [
    {"role": "user", "content": "Hello, Claude"}
  ]
}'
```

Response:

```
{
  "id": "msg_01XFDUDYJgAACzvnptvVoYEL",
  "type": "message",
  "role": "assistant",
  "content": [
    {
      "type": "text",
      "text": "Hello! How can I assist you today?"
    }
  ],
  "model": "claude-sonnet-4-5",
  "stop_reason": "end_turn",
  "usage": {
    "input_tokens": 12,
    "output_tokens": 8
  }
}
```

For complete examples and tutorials, see [Get started](#) and [Working with Messages](#).

Next Steps



in multiple languages

conversations, and best practices



Messages API Reference

Complete API specification: parameters, responses, and error codes



Client SDKs

Installation guides for Python, TypeScript, Java, Go, C#, Ruby, and PHP



Features overview

Explore capabilities: caching, vision, tool use, streaming, and more



Rate limits

Usage tiers, spend limits, and rate limiting with token bucket algorithm

[Security and compliance](#)

[Transparency](#)